

# Image Captioning with LSTMs

```
In [12]: # As usual, a bit of setup
from __future__ import print_function
import time, os, json
import numpy as np
import matplotlib.pyplot as plt

from deeplearning.gradient_check import eval_numerical_gradient, eval_
numerical_gradient_array
from deeplearning.rnn_layers import *
from deeplearning.captioning_solver import CaptioningSolver
from deeplearning.classifiers.rnn import CaptioningRNN
from deeplearning.coco_utils import load_coco_data, sample_coco_minib
atch, decode_captions
from deeplearning.image_utils import image_from_url

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of pl
ots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-module
s-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.ab
s(y))))
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

## Load MS-COCO data

As in the previous notebook, we will use the Microsoft COCO dataset for captioning.

```
In [13]: # Load COCO data from disk; this returns a dictionary
# We'll work with dimensionality-reduced features for this notebook
data = load_coco_data(pca_features=True)

# Print out all the keys and values from the data dictionary
for k, v in data.items():
    if type(v) == np.ndarray:
        print(k, type(v), v.shape, v.dtype)
    else:
        print(k, type(v), len(v))

train_captions <class 'numpy.ndarray'> (400135, 17) int32
train_image_idxs <class 'numpy.ndarray'> (400135,) int32
val_captions <class 'numpy.ndarray'> (195954, 17) int32
val_image_idxs <class 'numpy.ndarray'> (195954,) int32
train_features <class 'numpy.ndarray'> (82783, 512) float32
val_features <class 'numpy.ndarray'> (40504, 512) float32
idx_to_word <class 'list'> 1004
word_to_idx <class 'dict'> 1004
train_urls <class 'numpy.ndarray'> (82783,) <U63
val_urls <class 'numpy.ndarray'> (40504,) <U63
```

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

## Overfit LSTM captioning model

```
In [19]: np.random.seed(231)

small_data = load_coco_data(max_train=50)


small_lstm_model = CaptioningRNN(
    cell_type='lstm',
    word_to_idx=data['word_to_idx'],
    input_dim=data['train_features'].shape[1],
    hidden_dim=512,
    wordvec_dim=256,
    dtype=np.float32,
)

small_lstm_solver = CaptioningSolver(small_lstm_model, small_data,
    update_rule='adam',
    num_epochs=100,
    batch_size=50,
    optim_config={
        'learning_rate': 5e-3,
    },
    lr_decay=0.995,
    verbose=True, print_every=10,
)

small_lstm_solver.train()

# Plot the training losses
plt.plot(small_lstm_solver.loss_history)
plt.xlabel('Iteration')
plt.ylabel('Loss')
plt.title('Training loss history')
plt.show()
```

```
(Iteration 1 / 100) loss: 80.600535  
(Iteration 11 / 100) loss: 40.194296  
(Iteration 21 / 100) loss: 20.591769  
(Iteration 31 / 100) loss: 8.602131  
(Iteration 41 / 100) loss: 3.736282  
(Iteration 51 / 100) loss: 0.991312  
(Iteration 61 / 100) loss: 0.463356  
(Iteration 71 / 100) loss: 0.227835  
(Iteration 81 / 100) loss: 0.142447  
(Iteration 91 / 100) loss: 0.098880
```



## LSTM test-time sampling

```
In [20]: for split in ['train', 'val']:
          minibatch = sample_coco_minibatch(small_data, split=split, batch_size=2)
          gt_captions, features, urls = minibatch
          gt_captions = decode_captions(gt_captions, data['idx_to_word'])

          sample_captions = small_lstm_model.sample(features)
          sample_captions = decode_captions(sample_captions, data['idx_to_word'])

          for gt_caption, sample_caption, url in zip(gt_captions, sample_captions, urls):
              plt.imshow(image_from_url(url))
              plt.title('%s\n%s\nGT:%s' % (split, sample_caption, gt_caption))
              plt.axis('off')
              plt.show()
```

train  
a bedroom with a striped <UNK> and red walls <END>  
GT:<START> a bedroom with a striped <UNK> and red walls <END>



train  
a truck that is <UNK> <UNK> and a school bus <END>  
GT:<START> a truck that is <UNK> <UNK> and a school bus <END>



val  
a boy is sitting on the cement in his hand <END>  
GT:<START> a man is standing near a <UNK> bed <END>





val  
half a <UNK> <UNK> on the <UNK> <END>  
GT:<START> a <UNK> in a small boat leaves <UNK> in the water <END>



Train a good captioning model!

```

In [21]: import nltk

def BLEU_score(gt_caption, sample_caption):
    """
    gt_caption: string, ground-truth caption
    sample_caption: string, your model's predicted caption
    Returns unigram BLEU score.
    """
    reference = [x for x in gt_caption.split(' ')
                  if ('<END>' not in x and '<START>' not in x and '<UNK>' not in x)]
    hypothesis = [x for x in sample_caption.split(' ')
                  if ('<END>' not in x and '<START>' not in x and '<UNK>' not in x)]
    BLEUScore = nltk.translate.bleu_score.sentence_bleu([reference], hypothesis, weights = [1])
    return BLEUScore

def evaluate_model(model):
    """
    model: CaptioningRNN model
    Prints unigram BLEU score averaged over 1000 training and val examples.
    """
    BLEUScores = {}
    for split in ['train', 'val']:
        minibatch = sample_coco_minibatch(data, split=split, batch_size=1000)
        gt_captions, features, urls = minibatch
        gt_captions = decode_captions(gt_captions, data['idx_to_word'])

        sample_captions = model.sample(features)
        sample_captions = decode_captions(sample_captions, data['idx_to_word'])

        total_score = 0.0
        for gt_caption, sample_caption, url in zip(gt_captions, sample_captions, urls):
            total_score += BLEU_score(gt_caption, sample_caption)

        BLEUScores[split] = total_score / len(sample_captions)

    for split in BLEUScores:
        print('Average BLEU score for %s: %f' % (split, BLEUScores[split]))
    # smaller_data=load_coco_data(max_train=10000)
    # lstm_model = CaptioningRNN(
    #     cell_type='lstm',
    #     word_to_idx=data['word_to_idx'],
    #     input_dim=data['train_features'].shape[1],
    #     hidden_dim=512,
    #     wordvec_dim=256,
    #     dtype=np.float32,
    # )

```

```
# lstm_solver = CaptioningSolver(lstm_model, smaller_data,
#                               update_rule='adam',
#                               num_epochs=15,
#                               batch_size=50,
#                               optim_config={
#                               'learning_rate': 10e-3,
#                               },
#                               lr_decay=0.8,
#                               verbose=True, print_every=100,
#                               )

# lstm_solver.train()

# evaluate_model(lstm_model)
```

In [ ]:

```
In [ ]: smaller_data=load_coco_data()
lstm_model = CaptioningRNN(
    cell_type='lstm',
    word_to_idx=data['word_to_idx'],
    input_dim=data['train_features'].shape[1],
    hidden_dim=512,
    wordvec_dim=256,
    dtype=np.float32,
)

lstm_solver = CaptioningSolver(lstm_model, smaller_data,
    update_rule='adam',
    num_epochs=15,
    batch_size=100,
    optim_config={
        'learning_rate': 8e-3,
    },
    lr_decay=0.4,
    verbose=True, print_every=100,
)

lstm_solver.train()

evaluate_model(lstm_model)

(Iteration 1 / 60015) loss: 74.972106
```

In [ ]: