

In [ ]:

```

In [11]: # As usual, a bit of setup

from __future__ import print_function
import time, os, json
import numpy as np
import matplotlib.pyplot as plt

from deeplearning.gradient_check import eval_numerical_gradient, eval_
numerical_gradient_array
from deeplearning.rnn_layers import *
from deeplearning.captioning_solver import CaptioningSolver
from deeplearning.classifiers.rnn import CaptioningRNN
from deeplearning.coco_utils import load_coco_data, sample_coco_minib
atch, decode_captions
from deeplearning.image_utils import image_from_url

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of pl
ots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-module
s-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.ab
s(y))))

```

## Install h5py

The COCO dataset we will be using is stored in HDF5 format. To load HDF5 files, we will need to install the h5py Python package. From the command line, run:

```
pip install h5py
```

If you receive a permissions error, you may need to run the command as root:

```
sudo pip install h5py
```

You can also run commands directly from the Jupyter notebook by prefixing the command with the "!" character:

In [10]: !pip install h5py

Requirement already satisfied: h5py in ./env/lib/python3.6/site-packages (2.7.0)

Requirement already satisfied: six in ./env/lib/python3.6/site-packages (from h5py) (1.10.0)

Requirement already satisfied: numpy>=1.7 in ./env/lib/python3.6/site-packages (from h5py) (1.12.1)

```
In [12]: # Load COCO data from disk; this returns a dictionary
# We'll work with dimensionality-reduced features for this notebook,
# but feel
# free to experiment with the original features by changing the flag
# below.
data = load_coco_data(pca_features=True)

# Print out all the keys and values from the data dictionary
for k, v in data.items():
    if type(v) == np.ndarray:
        print(k, type(v), v.shape, v.dtype)
    else:
        print(k, type(v), len(v))
```

```
train_captions <class 'numpy.ndarray'> (400135, 17) int32
train_image_idxs <class 'numpy.ndarray'> (400135,) int32
val_captions <class 'numpy.ndarray'> (195954, 17) int32
val_image_idxs <class 'numpy.ndarray'> (195954,) int32
train_features <class 'numpy.ndarray'> (82783, 512) float32
val_features <class 'numpy.ndarray'> (40504, 512) float32
idx_to_word <class 'list'> 1004
word_to_idx <class 'dict'> 1004
train_urls <class 'numpy.ndarray'> (82783,) <U63
val_urls <class 'numpy.ndarray'> (40504,) <U63
```

```
In [13]: # Sample a minibatch and show the images and captions
batch_size = 3

captions, features, urls = sample_coco_minibatch(data, batch_size=batch_size)
for i, (caption, url) in enumerate(zip(captions, urls)):
    plt.imshow(image_from_url(url))
    plt.axis('off')
    caption_str = decode_captions(caption, data['idx_to_word'])
    plt.title(caption_str)
    plt.show()
```

<START> a woman wearing a short <UNK> kneeling on a tennis court <END>



<START> a tour bus driving through an empty intersection <END>



<START> a person skiing on skis on a snow covered hill <END>



In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

```
In [24]: np.random.seed(231)

small_data = load_coco_data(max_train=50)

small_rnn_model = CaptioningRNN(
    cell_type='rnn',
    word_to_idx=data['word_to_idx'],
    input_dim=data['train_features'].shape[1],
    hidden_dim=512,
    wordvec_dim=256,
)


small_rnn_solver = CaptioningSolver(small_rnn_model, small_data,
    update_rule='adam',
    num_epochs=50,
    batch_size=25,
    optim_config={
        'learning_rate': 5e-3,
    },
    lr_decay=0.95,
    verbose=True, print_every=10,
)

small_rnn_solver.train()

# Plot the training losses
plt.plot(small_rnn_solver.loss_history)
plt.xlabel('Iteration')
plt.ylabel('Loss')
plt.title('Training loss history')
plt.show()
```



```
(Iteration 1 / 100) loss: 76.913487  
(Iteration 11 / 100) loss: 21.062919  
(Iteration 21 / 100) loss: 4.016281  
(Iteration 31 / 100) loss: 0.566975  
(Iteration 41 / 100) loss: 0.239450  
(Iteration 51 / 100) loss: 0.161949  
(Iteration 61 / 100) loss: 0.111534  
(Iteration 71 / 100) loss: 0.097569  
(Iteration 81 / 100) loss: 0.099074  
(Iteration 91 / 100) loss: 0.073962
```



## Test-time sampling



```
In [37]: for split in ['train', 'val']:
          minibatch = sample_coco_minibatch(small_data, split=split, batch_size=2)
          gt_captions, features, urls = minibatch
          gt_captions = decode_captions(gt_captions, data['idx_to_word'])

          sample_captions = small_rnn_model.sample(features)
          sample_captions = decode_captions(sample_captions, data['idx_to_word'])

          for gt_caption, sample_caption, url in zip(gt_captions, sample_captions, urls):
              plt.imshow(image_from_url(url))
              plt.title('%s\n%s\nGT:%s' % (split, sample_caption, gt_caption))
              plt.axis('off')
              plt.show()
```

train  
a closeup of a woman from the <UNK> to <UNK> <END>  
GT:<START> a closeup of a woman from the <UNK> to <UNK> <END>



train  
a man standing on the side of a road with bags of luggage <END>  
GT:<START> a man standing on the side of a road with bags of luggage <END>





val  
sitting atop the pile of skiing gear in room <UNK> with wood <END>  
GT:<START> a man looks <UNK> and holds up his kite <END>



val  
a dog is on the <END>  
GT:<START> a white bed in the middle of a home room <END>



In [ ]:

In [ ]: