

# Style Transfer

In this notebook we will implement the style transfer technique from "[Image Style Transfer Using Convolutional Neural Networks](http://www.cv-foundation.org/openaccess/content_cvpr_2016/papers/Gatys_Image_Style_Transfer_CVPR_2016_paper.pdf)" (Gatys et al., CVPR 2015) ([http://www.cv-foundation.org/openaccess/content\\_cvpr\\_2016/papers/Gatys\\_Image\\_Style\\_Transfer\\_CVPR\\_2016\\_paper.pdf](http://www.cv-foundation.org/openaccess/content_cvpr_2016/papers/Gatys_Image_Style_Transfer_CVPR_2016_paper.pdf)).

The general idea is to take two images, and produce a new image that reflects the content of one but the artistic "style" of the other. We will do this by first formulating a loss function that matches the content and style of each respective image in the feature space of a deep network, and then performing gradient descent on the pixels of the image itself.

The deep network we use as a feature extractor is [SqueezeNet](https://arxiv.org/abs/1602.07360) (<https://arxiv.org/abs/1602.07360>), a small model that has been trained on ImageNet. You could use any network, but we chose SqueezeNet here for its small size and efficiency.

Here's an example of the images you'll be able to produce by the end of this notebook:



We will then use this to create a video style transferer and then use it to automatic download videos and styles from the internet and upload them to my [Youtube Channel](https://www.youtube.com/channel/UC2RKwvGB9hrVYQtmPNtCLkw/featured) (<https://www.youtube.com/channel/UC2RKwvGB9hrVYQtmPNtCLkw/featured>)

## Setup

In [ ]:

```
In [1]: import torch
import torch.nn as nn
from torch.autograd import Variable
import torchvision
import torchvision.transforms as T
import PIL

import numpy as np

import cv2
import os
from collections import namedtuple
import matplotlib.pyplot as plt

from deeplearning.image_utils import SQUEEZENET_MEAN, SQUEEZENET_STD
%matplotlib inline
```

These are helper functions to deal with images, since we're dealing with real JPEGs, not CIFAR-10 data.

In [ ]:

```

In [2]: def preprocess(img, size=512):
        transform = T.Compose([
            T.Scale((size,size)),
            T.ToTensor(),
            T.Normalize(mean=SQUEEZENET_MEAN.tolist(),
                        std=SQUEEZENET_STD.tolist()),
            T.Lambda(lambda x: x[None]),
        ])
        return transform(img)

def deprocess(img):
    transform = T.Compose([
        T.Lambda(lambda x: x[0]),
        T.Normalize(mean=[0, 0, 0], std=[1.0 / s for s in SQUEEZENET_
STD.tolist()]),
        T.Normalize(mean=[-m for m in SQUEEZENET_MEAN.tolist()], std=
[1, 1, 1]),
        T.Lambda(rescale),
        T.ToPILImage(),
    ])
    return transform(img)

def rescale(x):
    low, high = x.min(), x.max()
    x_rescaled = (x - low) / (high - low)
    return x_rescaled

def rel_error(x,y):
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.ab
s(y))))

def features_from_img(imgpath, imgsize):
    img = preprocess(PIL.Image.open(imgpath), size=imgsize)
    img_var = Variable(img.type(dtype))
    device = torch.device('cuda')
    return extract_features(img_var.to(device), cnn), img_var

# Older versions of scipy.misc.imresize yield different results
# from newer versions, so we check to make sure scipy is up to date.
def check_scipy():
    import scipy
    vnum = int(scipy.__version__.split('.')[1])

    assert vnum >= 16, "You must install SciPy >= 0.16.0 to complete
this notebook."

#check_scipy()

answers = np.load('style-transfer-checks.npz')

```

We need to set the dtype to select either the CPU or the GPU

```
In [3]: #dtype = torch.FloatTensor
        #comment above or below depending if you're on a machine with a GPU set up for PyTorch!
        dtype = torch.cuda.FloatTensor
```

```
In [4]: # Load the pre-trained SqueezeNet model.
        cnn = torchvision.models.squeezenet1_1(pretrained=True).features
        cnn.type(dtype)

        # We don't want to train the model any further, so we don't want PyTorch to waste computation
        # computing gradients on parameters we're never going to update.
        for param in cnn.parameters():
            param.requires_grad = False

        #helper code which takes an image, a model (cnn), and returns a list of
        # feature maps, one per layer.
        def extract_features(x, cnn):
            """
            Use the CNN to extract features from the input image x.

            Inputs:
            - x: A PyTorch Variable of shape (N, C, H, W) holding a minibatch of images that
              will be fed to the CNN.
            - cnn: A PyTorch model that we will use to extract features.

            Returns:
            - features: A list of feature for the input images x extracted using the cnn model.
              features[i] is a PyTorch Variable of shape (N, C_i, H_i, W_i); recall that features
              from different layers of the network may have different numbers of channels (C_i) and
              spatial dimensions (H_i, W_i).
            """
            features = []
            device = torch.device('cuda')
            cnn = cnn.to(device)
            prev_feat = x
            for i, module in enumerate(cnn._modules.values()):
                next_feat = module(prev_feat)
                features.append(next_feat)
                prev_feat = next_feat
            return features
```

## Computing Loss

We're going to compute the three components of our loss function now. The loss function is a weighted sum of three terms: content loss + style loss + total variation loss.

## Content loss

We can generate an image that reflects the content of one image and the style of another by incorporating both in our loss function. We want to penalize deviations from the content of the content image and deviations from the style of the style image. We can then use this hybrid loss function to perform gradient descent **not on the parameters** of the model, but instead **on the pixel values** of our original image.

Let's first write the content loss function. Content loss measures how much the feature map of the generated image differs from the feature map of the source image. We only care about the content representation of one layer of the network (say, layer  $\ell$ ), that has feature maps  $A^\ell \in \mathbb{R}^{1 \times C_\ell \times H_\ell \times W_\ell}$ .  $C_\ell$  is the number of filters/channels in layer  $\ell$ ,  $H_\ell$  and  $W_\ell$  are the height and width. We will work with reshaped versions of these feature maps that combine all spatial positions into one dimension. Let  $F^\ell \in \mathbb{R}^{N_\ell \times M_\ell}$  be the feature map for the current image and  $P^\ell \in \mathbb{R}^{N_\ell \times M_\ell}$  be the feature map for the content source image where  $M_\ell = H_\ell \times W_\ell$  is the number of elements in each feature map. Each row of  $F^\ell$  or  $P^\ell$  represents the vectorized activations of a particular filter, convolved over all positions of the image. Finally, let  $w_c$  be the weight of the content loss term in the loss function.

Then the content loss is given by:

$$L_c = w_c \times \sum_{i,j} (F_{ij}^\ell - P_{ij}^\ell)^2$$

```
In [5]: def content_loss(content_weight, content_current, content_original):
        """
        Compute the content loss for style transfer.

        Inputs:
        - content_weight: Scalar giving the weighting for the content loss.
        - content_current: features of the current image; this is a PyTorch Tensor of shape
            (1, C_l, H_l, W_l).
        - content_target: features of the content image, Tensor with shape (1, C_l, H_l, W_l).

        Returns:
        - scalar content loss
        """
        # sumy=0
        N,C,H,W=content_current.shape
        # for k in range(C):
        #     for i in range(H):
        #         for j in range(W):
        #             sumy+=(content_current[0][k][i][j]-content_original[0][k][i][j])**2

        sumy=torch.nn.functional.mse_loss(content_current,content_original)
        return N*C*H*W*content_weight*sumy
```

```
In [ ]:
```

## Style loss

Now we can tackle the style loss. For a given layer  $\ell$ , the style loss is defined as follows:

First, compute the Gram matrix  $G$  which represents the correlations between the responses of each filter, where  $F$  is as above. The Gram matrix is an approximation to the covariance matrix -- we want the activation statistics of our generated image to match the activation statistics of our style image, and matching the (approximate) covariance is one way to do that. There are a variety of ways you could do this, but the Gram matrix is nice because it's easy to compute and in practice shows good results.

Given a feature map  $F^\ell$  of shape  $(1, C_\ell, M_\ell)$ , the Gram matrix has shape  $(1, C_\ell, C_\ell)$  and its elements are given by:

$$G_{ij}^\ell = \sum_k F_{ik}^\ell F_{jk}^\ell$$

Assuming  $G^\ell$  is the Gram matrix from the feature map of the current image,  $A^\ell$  is the Gram Matrix from the feature map of the source style image, and  $w_\ell$  a scalar weight term, then the style loss for the layer  $\ell$  is simply the weighted Euclidean distance between the two Gram matrices:

$$L_s^\ell = w_\ell \sum_{i,j} \left( G_{ij}^\ell - A_{ij}^\ell \right)^2$$

In practice we usually compute the style loss at a set of layers  $\mathcal{L}$  rather than just a single layer  $\ell$ ; then the total style loss is the sum of style losses at each layer:

$$L_s = \sum_{\ell \in \mathcal{L}} L_s^\ell$$

We have implemented the Gram matrix computation below:

```

In [6]: def gram_matrix(features, normalize=True):
        """
        Compute the Gram matrix from features.

        Inputs:
        - features: PyTorch Variable of shape (N, C, H, W) giving feature
        s for
          a batch of N images.
        - normalize: optional, whether to normalize the Gram matrix
          If True, divide the Gram matrix by the number of neurons (H *
        W * C)

        Returns:
        - gram: PyTorch Variable of shape (N, C, C) giving the
          (optionally normalized) Gram matrices for the N input images.
        """
        a, b, c, d = features.size()
        features = features.view(a * b, c * d) # resise F_XL into \hat F
        _XL

        G = torch.mm(features, features.t()) # compute the gram product

        # we 'normalize' the values of the gram matrix
        # by dividing by the number of element in each feature maps.
        if normalize:
            return G.div(a * b * c * d)
        return G

        # trans=features.transpose(1,2)
        # print(trans.shape)
        # gram=trans.dot(features)

```

In [ ]:

In [ ]:

Next, we implement the style loss:

```

In [7]: # Now put it together in the style_loss function...
def style_loss(feats, style_layers, style_targets, style_weights):
    """
    Computes the style loss at a set of layers.

    Inputs:
    - feats: list of the features at every layer of the current image, as produced by the extract_features function.
    - style_layers: List of layer indices into feats giving the layers to include in the style loss.
    - style_targets: List of the same length as style_layers, where style_targets[i] is a PyTorch Variable giving the Gram matrix the source style image computed at layer style_layers[i].
    - style_weights: List of the same length as style_layers, where style_weights[i] is a scalar giving the weight for the style loss at layer style_layers[i].

    Returns:
    - style_loss: A PyTorch Variable holding a scalar giving the style loss.
    """
    # for l in range(len(style_layers)):
    #     G=gram_matrix(feats[style_layers[l]])

    #     C,C2=G.size()
    #     summy=0

    #     for i in range(C):
    #         for j in range(C2):
    #             #print((i,j))
    #             summy+=(G[i][j]-style_targets[l][i][j])**2
    #     losses.append(style_weights[l]*summy)
    #     return sum(losses)
    losses=[]
    for l in range(len(style_layers)):
        G=gram_matrix(feats[style_layers[l]])
        C,_=G.shape
        #print(G,style_targets[l],)
        #print(G.shape,style_targets[l].shape)
        losses.append(style_weights[l]*(C**2*torch.nn.functional.mse_loss(G, style_targets[l])))
    #print(losses)
    return sum(losses)

```

In [ ]:

In [ ]:



## Total-variation regularization

It turns out that it's helpful to also encourage smoothness in the image. We can do this by adding another term to our loss that penalizes wiggles or "total variation" in the pixel values.

You can compute the "total variation" as the sum of the squares of differences in the pixel values for all pairs of pixels that are next to each other (horizontally or vertically). Here we sum the total-variation regularization for each of the 3 input channels (RGB), and weight the total summed loss by the total variation weight,  $w_t$ :

$$L_{tv} = w_t \times \sum_{c=1}^3 \sum_{i=1}^{H-1} \sum_{j=1}^{W-1} ((x_{i,j+1,c} - x_{i,j,c})^2 + (x_{i+1,j,c} - x_{i,j,c})^2)$$

In the next cell, we fill in the definition for the TV loss term.

```
In [8]: def tv_loss(img, tv_weight):
        """
        Compute total variation loss.

        Inputs:
        - img: PyTorch Variable of shape (1, 3, H, W) holding an input im
          age.
        - tv_weight: Scalar giving the weight w_t to use for the TV loss.

        Returns:
        - loss: PyTorch Variable holding a scalar giving the total variat
          ion loss
          for img weighted by tv_weight.
        """
        # Your implementation should be vectorized and not require any lo
        ops!
        diff_i = torch.sum(torch.square((img[:, :, :, 1:] - img[:, :, :,
          :-1])))
        diff_j = torch.sum(torch.square(torch.abs(img[:, :, 1:, :] - img
         [:, :, :-1, :])))
        tv_loss = tv_weight*(diff_i + diff_j)
        return tv_loss
```

In [ ]:

In [ ]:

Now we're ready to string it all together:

```

In [9]: def style_transfer(content_image, style_image, image_size, style_size
, content_layer, content_weight,
            style_layers, style_weights, tv_weight, init_rando
m = False, quiet=False):
    """
    Run style transfer!

    Inputs:
    - content_image: filename of content image
    - style_image: filename of style image
    - image_size: size of smallest image dimension (used for content
loss and generated image)
    - style_size: size of smallest style image dimension
    - content_layer: layer to use for content loss
    - content_weight: weighting on content loss
    - style_layers: list of layers to use for style loss
    - style_weights: list of weights to use for each layer in style_l
ayers
    - tv_weight: weight of total variation regularization term
    - init_random: initialize the starting image to uniform random no
ise
    """

    device = torch.device('cuda')
    # Extract features for the content image
    content_img = preprocess(PIL.Image.open(content_image), size=imag
e_size)
    content_img_var = Variable(content_img.type(dtype)).to(device)
    feats = extract_features(content_img_var, cnn)
    content_target = feats[content_layer].clone().to(device)

    # Extract features for the style image
    style_img = preprocess(PIL.Image.open(style_image), size=style_si
ze).to(device)
    style_img_var = Variable(style_img.type(dtype)).to(device)
    feats = extract_features(style_img_var, cnn)
    style_targets = []
    for idx in style_layers:
        style_targets.append(gram_matrix(feats[idx].clone()))

    # Initialize output image to content image or noise
    if init_random:
        img = torch.Tensor(content_img.size()).uniform_(0, 1)
    else:
        img = content_img.clone().type(dtype).to(device)
    # We do want the gradient computed on our image!
    img_var = Variable(img, requires_grad=True).to(device)

    # Set up optimization hyperparameters
    initial_lr = 3.0
    decayed_lr = 0.1
    decay_lr_at = 180

    # Note that we are optimizing the pixel values of the image by pa
ssing

```

```

# in the img_var Torch variable, whose requires_grad flag is set
to True
optimizer = torch.optim.Adam([img_var], lr=initial_lr)
if not quiet:
    f, axarr = plt.subplots(1,2)
    axarr[0].axis('off')
    axarr[1].axis('off')
    axarr[0].set_title('Content Source Img.')
    axarr[1].set_title('Style Source Img.')
    axarr[0].imshow(deprocess(content_img.cpu()))
    axarr[1].imshow(deprocess(style_img.cpu()))
    plt.show()
    plt.figure()

for t in range(200):
    if t < 190:
        img.clamp_(-1.5, 1.5)
        optimizer.zero_grad()

        feats = extract_features(img_var, cnn)
        #print(t)
        # Compute loss
        c_loss = content_loss(content_weight, feats[content_layer], c
content_target)
        #print(t)
        s_loss = style_loss(feats, style_layers, style_targets, style
_weights)
        #print(t)
        t_loss = tv_loss(img_var, tv_weight)
        #print(t)
        loss = c_loss + s_loss + t_loss

        loss.backward()
        #print(t)
        # Perform gradient descents on our image values
        if t == decay_lr_at:
            optimizer = torch.optim.Adam([img_var], lr=decayed_lr)
            optimizer.step()

        if t % 100 == 0:
            if not quiet:
                print('Iteration {}'.format(t))
                plt.axis('off')
                plt.imshow(deprocess(img.cpu()))
                plt.show()

return img

```

## Generate some pretty pictures!

- The `content_image` is the filename of content image.
- The `style_image` is the filename of style image.
- The `image_size` is the size of smallest image dimension of the content image (used for content loss and generated image).
- The `style_size` is the size of smallest style image dimension.
- The `content_layer` specifies which layer to use for content loss.
- The `content_weight` gives weighting on content loss in the overall loss function. Increasing the value of this parameter will make the final image look more realistic (closer to the original content).
- `style_layers` specifies a list of which layers to use for style loss.
- `style_weights` specifies a list of weights to use for each layer in `style_layers` (each of which will contribute a term to the overall style loss). We generally use higher weights for the earlier style layers because they describe more local/smaller scale features, which are more important to texture than features over larger receptive fields. In general, increasing these weights will make the resulting image look less like the original content and more distorted towards the appearance of the style image.
- `tv_weight` specifies the weighting of total variation regularization in the overall loss function. Increasing this value makes the resulting image look smoother and less jagged, at the cost of lower fidelity to style and content.

```

In [18]: def styletrans(im1,im2, show=False):
    content_size = 224
    style_size = 224
    params1 = {
        'content_image' : im1,
        'style_image' : im2,
        'image_size' : content_size,
        'style_size' : style_size,
        'content_layer' : 3,
        'content_weight' : 1e-1,
        'style_layers' : [1, 4, 6, 7],
        'style_weights' : [300000, 1000, 15, 3],
        'tv_weight' : 5e-2,
        'quiet': True
    }

    img = style_transfer(**params1)
    arr = deprocess(img.cpu())
    if show:
        plt.axis('off')
        plt.imshow(arr)
        plt.show()
    return cv2.cvtColor(np.asarray(arr), cv2.COLOR_RGB2BGR)

#Movie Maker
def styletrans_a_mp4(movie, style, outputmovie,s=-1):
    cam = cv2.VideoCapture(movie)
    frame_ct= int(cam.get(cv2.CAP_PROP_FRAME_COUNT))
    # initialize video writer
    fourcc = cv2.VideoWriter_fourcc('M','P','E','G')
    fps = int(cam.get(cv2.CAP_PROP_FPS))
    print("FPS:",fps)
    video_filename = outputmovie
    width = 224
    height = 224
    out = cv2.VideoWriter(video_filename, fourcc, fps, (width, height
))
    done_ct = 0
    print(str(done_ct)+"/"+str(int((frame_ct if s== -1 else s*fps))))
    while(True):
        # reading from frame
        ret,frame = cam.read()
        if ret:
            if s!=-1 and done_ct > int(s*fps):
                print("Stopping after " + str(s) + " seconds of foota
ge")
                out.release()
                return
            done_ct+=1
            name = 'xyzfdsf.jpg'
            cv2.imwrite(name, frame)
            if done_ct % int(fps/8) == 0:
                print(str(done_ct)+"/"+str(int((frame_ct if s== -1 els
e s*fps))))
            img = styletrans(name, style, show = (done_ct ==1))

```

```
        out.write(img)
        os.remove(name)
    else:
        out.release()
    return
```



```
arr = styletrans('styles/tubingen.jpg', 'styles/starry_night.jpg', show=True)  
print(arr.shape)
```

```
/home/rishirules/anaconda3/lib/python3.7/site-packages/torchvision/transforms/transforms.py:285: UserWarning: The use of the transforms.Scale transform is deprecated, please use transforms.Resize instead.  
  warnings.warn("The use of the transforms.Scale transform is deprecated, " +
```



(224, 224, 3)



```
In [12]: #test movie style transfer  
styletrans_a_mp4('styles/dog.mp4', 'styles/fire_demon_van_gogh.jpg', 'styles/fire_dog.mp4')
```

FPS: 59  
0/3537



7/3537  
14/3537  
21/3537  
28/3537  
35/3537  
42/3537  
49/3537  
56/3537  
63/3537  
70/3537  
77/3537  
84/3537  
91/3537  
98/3537  
105/3537  
112/3537  
119/3537  
126/3537  
133/3537  
140/3537  
147/3537  
154/3537  
161/3537  
168/3537  
175/3537  
182/3537  
189/3537  
196/3537  
203/3537  
210/3537  
217/3537  
224/3537  
231/3537  
238/3537  
245/3537  
252/3537  
259/3537  
266/3537  
273/3537  
280/3537  
287/3537  
294/3537  
301/3537  
308/3537  
315/3537  
322/3537  
329/3537  
336/3537  
343/3537  
350/3537  
357/3537  
364/3537  
371/3537  
378/3537  
385/3537  
392/3537  
399/3537

406/3537  
413/3537  
420/3537  
427/3537  
434/3537  
441/3537  
448/3537  
455/3537  
462/3537  
469/3537  
476/3537  
483/3537  
490/3537  
497/3537  
504/3537  
511/3537  
518/3537  
525/3537  
532/3537  
539/3537  
546/3537  
553/3537  
560/3537  
567/3537  
574/3537  
581/3537  
588/3537  
595/3537  
602/3537  
609/3537  
616/3537  
623/3537  
630/3537  
637/3537  
644/3537  
651/3537  
658/3537  
665/3537  
672/3537  
679/3537  
686/3537  
693/3537  
700/3537  
707/3537  
714/3537  
721/3537  
728/3537  
735/3537  
742/3537  
749/3537  
756/3537  
763/3537  
770/3537  
777/3537  
784/3537  
791/3537  
798/3537

805/3537  
812/3537  
819/3537  
826/3537  
833/3537  
840/3537  
847/3537  
854/3537  
861/3537  
868/3537  
875/3537  
882/3537  
889/3537  
896/3537  
903/3537  
910/3537  
917/3537  
924/3537  
931/3537  
938/3537  
945/3537  
952/3537  
959/3537  
966/3537  
973/3537  
980/3537  
987/3537  
994/3537  
1001/3537  
1008/3537  
1015/3537  
1022/3537  
1029/3537  
1036/3537  
1043/3537  
1050/3537  
1057/3537  
1064/3537  
1071/3537  
1078/3537  
1085/3537  
1092/3537  
1099/3537  
1106/3537  
1113/3537  
1120/3537  
1127/3537  
1134/3537  
1141/3537  
1148/3537  
1155/3537  
1162/3537  
1169/3537  
1176/3537  
1183/3537  
1190/3537  
1197/3537

1204/3537  
1211/3537  
1218/3537  
1225/3537  
1232/3537  
1239/3537  
1246/3537  
1253/3537  
1260/3537  
1267/3537  
1274/3537  
1281/3537  
1288/3537  
1295/3537  
1302/3537  
1309/3537  
1316/3537  
1323/3537  
1330/3537  
1337/3537  
1344/3537  
1351/3537  
1358/3537  
1365/3537  
1372/3537  
1379/3537  
1386/3537  
1393/3537  
1400/3537  
1407/3537  
1414/3537  
1421/3537  
1428/3537  
1435/3537  
1442/3537  
1449/3537  
1456/3537  
1463/3537  
1470/3537  
1477/3537  
1484/3537  
1491/3537  
1498/3537  
1505/3537  
1512/3537  
1519/3537  
1526/3537  
1533/3537  
1540/3537  
1547/3537  
1554/3537  
1561/3537  
1568/3537  
1575/3537  
1582/3537  
1589/3537  
1596/3537

1603/3537  
1610/3537  
1617/3537  
1624/3537  
1631/3537  
1638/3537  
1645/3537  
1652/3537  
1659/3537  
1666/3537  
1673/3537  
1680/3537  
1687/3537  
1694/3537  
1701/3537  
1708/3537  
1715/3537  
1722/3537  
1729/3537  
1736/3537  
1743/3537  
1750/3537  
1757/3537  
1764/3537  
1771/3537  
1778/3537  
1785/3537  
1792/3537  
1799/3537  
1806/3537  
1813/3537  
1820/3537  
1827/3537  
1834/3537  
1841/3537  
1848/3537  
1855/3537  
1862/3537  
1869/3537  
1876/3537  
1883/3537  
1890/3537  
1897/3537  
1904/3537  
1911/3537  
1918/3537  
1925/3537  
1932/3537  
1939/3537  
1946/3537  
1953/3537  
1960/3537  
1967/3537  
1974/3537  
1981/3537  
1988/3537  
1995/3537

2002/3537  
2009/3537  
2016/3537  
2023/3537  
2030/3537  
2037/3537  
2044/3537  
2051/3537  
2058/3537  
2065/3537  
2072/3537  
2079/3537  
2086/3537  
2093/3537  
2100/3537  
2107/3537  
2114/3537  
2121/3537  
2128/3537  
2135/3537  
2142/3537  
2149/3537  
2156/3537  
2163/3537  
2170/3537  
2177/3537  
2184/3537  
2191/3537  
2198/3537  
2205/3537  
2212/3537  
2219/3537  
2226/3537  
2233/3537  
2240/3537  
2247/3537  
2254/3537  
2261/3537  
2268/3537  
2275/3537  
2282/3537  
2289/3537  
2296/3537  
2303/3537  
2310/3537  
2317/3537  
2324/3537  
2331/3537  
2338/3537  
2345/3537  
2352/3537  
2359/3537  
2366/3537  
2373/3537  
2380/3537  
2387/3537  
2394/3537



2401/3537  
2408/3537  
2415/3537  
2422/3537  
2429/3537  
2436/3537  
2443/3537  
2450/3537  
2457/3537  
2464/3537  
2471/3537  
2478/3537  
2485/3537  
2492/3537  
2499/3537  
2506/3537  
2513/3537  
2520/3537  
2527/3537  
2534/3537  
2541/3537  
2548/3537  
2555/3537  
2562/3537  
2569/3537  
2576/3537  
2583/3537  
2590/3537  
2597/3537  
2604/3537  
2611/3537  
2618/3537  
2625/3537  
2632/3537  
2639/3537  
2646/3537  
2653/3537  
2660/3537  
2667/3537  
2674/3537  
2681/3537  
2688/3537  
2695/3537  
2702/3537  
2709/3537  
2716/3537  
2723/3537  
2730/3537  
2737/3537  
2744/3537  
2751/3537  
2758/3537  
2765/3537  
2772/3537  
2779/3537  
2786/3537  
2793/3537

2800/3537  
2807/3537  
2814/3537  
2821/3537  
2828/3537  
2835/3537  
2842/3537  
2849/3537  
2856/3537  
2863/3537  
2870/3537  
2877/3537  
2884/3537  
2891/3537  
2898/3537  
2905/3537  
2912/3537  
2919/3537  
2926/3537  
2933/3537  
2940/3537  
2947/3537  
2954/3537  
2961/3537  
2968/3537  
2975/3537  
2982/3537  
2989/3537  
2996/3537  
3003/3537  
3010/3537  
3017/3537  
3024/3537  
3031/3537  
3038/3537  
3045/3537  
3052/3537  
3059/3537  
3066/3537  
3073/3537  
3080/3537  
3087/3537  
3094/3537  
3101/3537  
3108/3537  
3115/3537  
3122/3537  
3129/3537  
3136/3537  
3143/3537  
3150/3537  
3157/3537  
3164/3537  
3171/3537  
3178/3537  
3185/3537  
3192/3537

3199/3537  
3206/3537  
3213/3537  
3220/3537  
3227/3537  
3234/3537  
3241/3537  
3248/3537  
3255/3537  
3262/3537  
3269/3537  
3276/3537  
3283/3537  
3290/3537  
3297/3537  
3304/3537  
3311/3537  
3318/3537  
3325/3537  
3332/3537  
3339/3537  
3346/3537  
3353/3537  
3360/3537  
3367/3537  
3374/3537  
3381/3537  
3388/3537  
3395/3537  
3402/3537  
3409/3537  
3416/3537  
3423/3537  
3430/3537  
3437/3537  
3444/3537  
3451/3537  
3458/3537  
3465/3537  
3472/3537  
3479/3537  
3486/3537  
3493/3537  
3500/3537  
3507/3537  
3514/3537  
3521/3537  
3528/3537  
3535/3537

1.00

0:00

## Other Stuff:

```
In [13]: # # Composition VII + Tubingen
# params1 = {
#     'content_image' : 'styles/tubingen.jpg',
#     'style_image' : 'styles/composition_vii.jpg',
#     'image_size' : 192,
#     'style_size' : 512,
#     'content_layer' : 3,
#     'content_weight' : 5e-2,
#     'style_layers' : (1, 4, 6, 7),
#     'style_weights' : (20000, 500, 12, 1),
#     'tv_weight' : 5e-2
# }

# style_transfer(**params1)
```

```
In [14]: !which ffmpeg

/usr/bin/ffmpeg
```

```
In [15]: # # Scream + Tubingen
# params2 = {
#     'content_image': 'styles/tubingen.jpg',
#     'style_image': 'styles/the_scream.jpg',
#     'image_size': 192,
#     'style_size': 224,
#     'content_layer': 3,
#     'content_weight': 3e-2,
#     'style_layers': [1, 4, 6, 7],
#     'style_weights': [200000, 800, 12, 1],
#     'tv_weight': 2e-2
# }

# style_transfer(**params2)
```

```
In [16]: # # Starry Night + Tübingen
# params3 = {
#     'content_image' : 'styles/tubingen.jpg',
#     'style_image' : 'styles/starry_night.jpg',
#     'image_size' : 192,
#     'style_size' : 192,
#     'content_layer' : 3,
#     'content_weight' : 6e-2,
#     'style_layers' : [1, 4, 6, 7],
#     'style_weights' : [300000, 1000, 15, 3],
#     'tv_weight' : 2e-2
# }

# style_transfer(**params3)
```

## Feature Inversion

The code written can do another cool thing. In an attempt to understand the types of features that convolutional networks learn to recognize, a recent paper [1] attempts to reconstruct an image from its feature representation. We can easily implement this idea using image gradients from the pretrained network, which is exactly what we did above (but with two different feature representations).

Now, if you set the style weights to all be 0 and initialize the starting image to random noise instead of the content source image, you'll reconstruct an image from the feature representation of the content source image. You're starting with total noise, but you should end up with something that looks quite a bit like your original image.

(Similarly, you could do "texture synthesis" from scratch if you set the content weight to 0 and initialize the starting image to random noise)

[1] Aravindh Mahendran, Andrea Vedaldi, "Understanding Deep Image Representations by Inverting them", CVPR 2015

Example below

```
In [17]: # # Feature Inversion -- Starry Night + Tubingen
# params_inv = {
#     'content_image' : 'styles/tubingen.jpg',
#     'style_image' : 'styles/starry_night.jpg',
#     'image_size' : 192,
#     'style_size' : 192,
#     'content_layer' : 3,
#     'content_weight' : 6e-2,
#     'style_layers' : [1, 4, 6, 7],
#     'style_weights' : [0, 0, 0, 0], # we discard any contributions
#     'tv_weight' : 2e-2,
#     'init_random': True # we want to initialize our image to be random
# }

# style_transfer(**params_inv)
```

In [ ]:

In [ ]: